

University of Central Florida

**STARS**

---

Data Science and Data Mining

---

January 2025

# Classification and Evaluation of Machine Learning Algorithms on the MNIST Dataset

Felix Yeboah

University of Central Florida, fe528610@ucf.edu



Part of the [Data Science Commons](#)

Find similar works at: <https://stars.library.ucf.edu/data-science-mining>

University of Central Florida Libraries <http://library.ucf.edu>

This Article is brought to you for free and open access by STARS. It has been accepted for inclusion in Data Science and Data Mining by an authorized administrator of STARS. For more information, please contact [STARS@ucf.edu](mailto:STARS@ucf.edu).

---

## STARS Citation

Yeboah, Felix, "Classification and Evaluation of Machine Learning Algorithms on the MNIST Dataset" (2025). *Data Science and Data Mining*. 27.

<https://stars.library.ucf.edu/data-science-mining/27>

Fig. 1. A 28x28 matrix of a single instance




Fig. 2. Distribution of Target labels

and dark pixels due to their dominance. In addition, Fig. 3 can be seen as the distribution of target variables. Nearly all the target labels have more than 6000 instances, with target label "1" having the most significant number of instances and target label "5" having the minimum number of instances. With regard to visualizing the distribution of the whole training set, a method called "t-SNE" was utilized. T-Distributed Stochastic Neighbor Embedding (t-SNE) visualizes high-dimensional data by giving each data point a location in a two or three-dimensional plane [6]. "t-SNE" attempts to ensure that each point has an equal number of neighbors by creating an idea of which other points are its "neighbors" for each point. After that, it attempts to embed them so that each point has an equal number of neighbors [7]. See Fig. 4. The next section presents two methods and evaluation metrics for classifying the MNIST dataset.


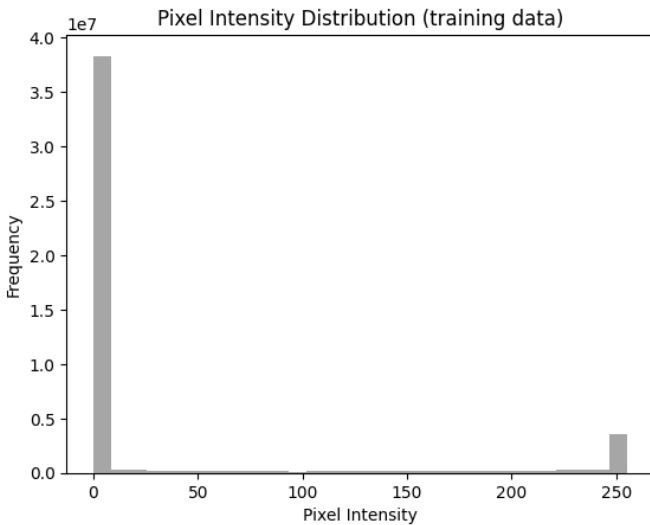


Fig. 4. Distribution of Training Set

### III. METHODOLOGY

#### A. *k*-Nearest Neighbor

The K-nearest neighbors (KNN) algorithm is a nonparametric supervised machine learning approach to classification. KNN tries to categorize the accurate class for the test sample by calculating the distance between the test sample and all the training samples. The algorithm looks for the highest number of classes or the average of the neighboring data points with the most common characteristics with the newly added data point. [8].

Consider a test sample  $\mathbf{x}$  that we wish to categorize into one of the  $K$  groups, we find the  $K$  observed data points that are closer to  $\mathbf{x}$ . The classification algorithm assigns  $\mathbf{x}$  to the population with the highest number of observed data points out of the  $K$ -nearest neighbors. When there is no majority vote the new instances are classified to one of the majority populations at random or left unclassified [9]. Algorithm 1 is pseudo-code for the KNN algorithm.

#### B. Naïve Bayes

Naive Bayes algorithm (NB) can broadly be defined as a family of machine learning classifiers that uses Bayes's theorem and conditional independence assumption to determine the probability that an instance belongs to a specific category. The algorithm is named naive because it assumes all the features are independent. This assumption is implausible in many real-world situations. Naive Bayes often produces competitive classification accuracy, although the assumption of independence is often violated in practice. Given a set of features  $X_1, X_2, X_3, \dots, X_n$ , the objective of classifying  $Y$ , is similar to maximizing

$$\arg \max_{y \in Y} P(Y = y | X_1, X_2, X_3, \dots, X_n)$$

Using Bayes' Theorem, we can rewrite this problem as

**Algorithm 1** Algorithm for k-Nearest Neighbor

**Input:** Given a set of training samples  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^{60000}$ , where:

- $X_i$  = ith feature instance,
- $y_i$  = ith target label,
- $K$  = Number of nearest neighbors,
- $\mathbf{x}$  = New instance.

**Output:** Class of test instance

- 1: **for** every  $\mathbf{x}$  **do**
- 2:   Compute the distance:  $d(\mathbf{x}, x) = \sqrt{\sum_{i=1}^N (\mathbf{x}_i - X_i)^2}$
- 3: **end for**
- 4: **return**  $\hat{y} = \arg \max_{c \in C} \sum_{i=1}^K C(y_{(i)} = c)$

$$\begin{aligned} P(Y = y | X_1, X_2, \dots, X_n) \\ = \frac{P(X_1, X_2, \dots, X_n | Y = y)P(Y = y)}{P(X_1, X_2, \dots, X_n)} \end{aligned} \quad (1)$$

Expanding numerator in “(1)”, we will have

$$\begin{aligned} P(Y = y | X_1, X_2, \dots, X_n) \\ = P(X_1, X_2, \dots, X_n | Y = y) \\ = P(X_1, X_2, \dots, X_n | Y = y)P(X_2, \dots, X_n | Y = y) \times \\ P(X_1 | X_2, \dots, X_n, Y = y)P(X_2 | X_3, \dots, X_n, Y = y) \\ \dots P(X_n | Y = y)P(Y = y) \end{aligned} \quad (2)$$

Modeling  $P(Y = y)P(X_1, X_2, \dots, X_n | Y = y)$  is difficult because there are too many parameters, and in most cases, we will run out of space and training data. Instead we assume that the features  $X_1, X_2, X_3, \dots, X_n$  are independent given the labels  $Y$ . We can reformulate “(2)” as

$$P(Y = y | X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | Y = y) \quad (3)$$

Therefore we have

$$\begin{aligned} P(Y = y | \mathbf{X}) \\ = \frac{P(Y = y)}{P(\mathbf{X})} \prod_{i=1}^n P(X_i | Y = y) \end{aligned} \quad (4)$$

“(4)” is equivalent to

$$\arg \max_{y \in Y} P(Y = y) \prod_{i=1}^n P(X_i | C = k) \quad (5)$$

Since  $P(\mathbf{X})$  is constant and does not depend on  $y$ , we drop it. This transformation makes the problem computable [10]. The methods described above were implemented in Python

using the Scikit-learn machine-learning library, and the results obtained from them are described in the next chapter.

## IV. RESULTS AND DISCUSSION

Having discussed the methods used in classifying the MNIST dataset, this section discusses, evaluates, and compares the results from these methods.

Table I shows the classification report after training the KNN and Naive Bayes algorithm on the 60000 training set and 10000 test set. Comparing KNN with  $k=3$  nearest neighbors to Gaussian naive Bayes, KNN had an accuracy rate of 97% and an error rate of 3%, and naive Bayes had an accuracy of 56% and an error rate of 44%. The poor performance of the naive Bayes algorithm can be attributed to the nonnormal nature of the pixel distribution. Reproducing the analysis with a distribution that closely models the MNIST data can improve the performance of the Naive Bayes model.

TABLE I  
PRECISION, RECALL AND F1-SCORE FOR KNN AND NAIVE BAYES

	KNN			Naive Bayes		
	Precision	Recall	F1-score	Precision	Recall	F1-score
0	0.96	0.99	0.98	0.79	0.89	0.84
1	0.95	1.00	0.98	0.85	0.95	0.90
2	0.98	0.96	0.97	0.90	0.26	0.40
3	0.96	0.97	0.97	0.71	0.35	0.47
4	0.98	0.96	0.97	0.88	0.17	0.29
5	0.97	0.97	0.97	0.55	0.05	0.09
6	0.98	0.99	0.98	0.65	0.93	0.77
7	0.96	0.96	0.96	0.88	0.27	0.42
8	0.99	0.94	0.96	0.28	0.67	0.40
9	0.96	0.95	0.95	0.37	0.95	0.53

<sup>a</sup>Error Rate KNN: 0.03.

<sup>b</sup>Accuracy Rate for KNN: 97%.

<sup>c</sup>Error Rate for Naive Bayes: 0.44.

<sup>d</sup>Accuracy Rate Naive Bayes: 56%.

The distribution of misclassified labels for KNN and naive Bayes is shown in Fig. 5 and Fig. 6, respectively, with KNN having fewer misclassified labels.




Fig. 5. Confusion Matrix for KNN

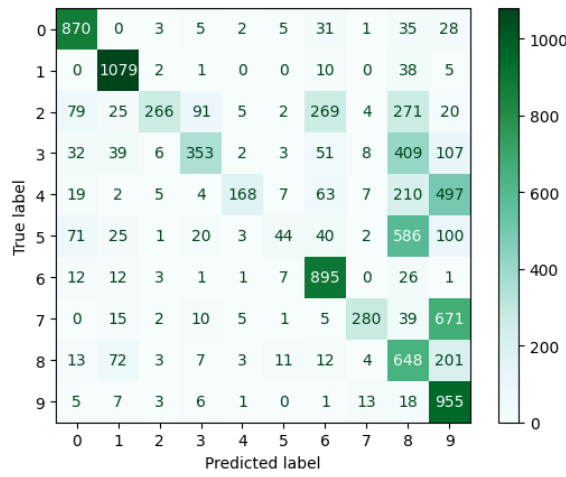


Fig. 6. Confusion Matrix for Naive Bayes

Receiver operating characteristic (ROC) curves of the KNN and Naive Bayes models are displayed in Fig. 7. The area under the ROC curve shows the probability that the model will perform better when given a new instance. Fig. 7 shows that the KNN model will likely perform better than the Naive Bayes model.

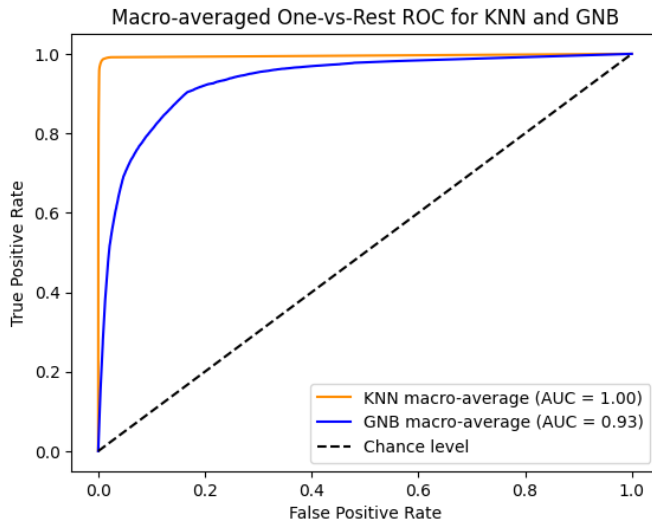


Fig. 7. ROC curve for KNN and Naive Bayes

## V. CONCLUSION

This project was undertaken to explore, evaluate, and compare machine learning algorithms on the widely-used MNIST dataset. Due to the simplicity of the MNIST dataset, it has widely been used to benchmark most machine-learning algorithms. Based on the evaluation metrics used in this study, the K-Nearest Neighbors (KNN) algorithm came up as the best model compared with Naive Bayes. Even though the KNN algorithm is computationally expensive, it can still provide good performance and predictions. One source of weakness in

this study that could have affected the Naive Bayes algorithm's performance is the pixel intensity distribution. For further studies, more information on pixel distribution would improve the performance of the Naive Bayes algorithm.

## REFERENCES

- [1] M. Shinozuka and B. Mansouri, "Synthetic aperture radar and remote sensing technologies for structural health monitoring of civil infrastructure systems," *Structural Health Monitoring of Civil Infrastructure Systems*, pp. 113–151, 2009, doi: <https://doi.org/10.1533/9781845696825.1.114>.
- [2] K. Huang, "Image Classification Using the Method of Convolutional Neural Networks," *2022 IEEE Conference on Telecommunications, Optics and Computer Science (TOCS)*, Dec. 2022, doi: <https://doi.org/10.1109/tocs56154.2022.10016070>.
- [3] W. Sun, "Image classification based on CNN with three different networks," *Applied and Computational Engineering*, vol. 15, no. 1, pp. 181–186, Oct. 2023, doi: <https://doi.org/10.54254/2755-2721/15/20230831>.
- [4] Yann LeCun et al., "Learning algorithms for classification: A comparison on handwritten digit recognition," 1995. Available: <https://api.semanticscholar.org/CorpusID:13411815>.
- [5] L. V. D. Maaten and G. E. Hinton, "Visualizing Data using t-SNE," 2008. <https://www.semanticscholar.org/paper/Visualizing-Data-using-t-SNE-Maaten-Hinton/1c46943103bd7b7a2c7be86859995a4144d1938b>
- [6] C. Olah, "Visualizing MNIST: An Exploration of Dimensionality Reduction - colah's blog," *colah.github.io*, Oct. 09, 2014. <https://colah.github.io/posts/2014-10-Visualizing-MNIST/> (accessed Sep. 22, 2024).
- [7] Rekha Gangula and Venkateswarlu B, "Exploring the Power and Practical Applications of K-Nearest Neighbours (KNN) in Machine Learning," vol. 2, no. 1, Feb. 2024, doi: <https://doi.org/10.69996/jcai.2024002>.
- [8] R. C. Neath and M. S. Johnson, "Discrimination and classification," P. Peterson, E. Baker, and B. McGaw, Eds., Third Edition. Elsevier, 2010, pp. 135–141. doi: <https://doi.org/10.1016/B978-0-08-044894-7.01312-9>.
- [9] Abedalmuhdi Almomany, W. R. Ayyad, and A. Jarrah, "Optimized implementation of an improved KNN classification algorithm using Intel FPGA platform: Covid-19 case study," *Journal of King Saud University - Computer and Information Sciences*, vol. 34, Art. no. 6, Part B, 2022, doi: <https://doi.org/10.1016/j.jksuci.2022.04.006>.
- [10] I. Rish, "An empirical study of the naive Bayes classifier," 2001. Available: <https://faculty.cc.gatech.edu/~isbell/reading/papers/Rish.pdf>

## VI. APPENDIX

```
1 # %% [markdown]
2 # # Importing libraries and Dataset
3
4 # %%
5
6 import pandas as pd
7 import matplotlib.pyplot as plt
8 import numpy as np
9 from sklearn.manifold import TSNE
10 %matplotlib inline
11 from keras.datasets import mnist
12 (x_train, y_train), (x_test, y_test) = mnist.
    load_data()
13
14 # %% [markdown]
15 # # Splitting, normalizing, reshaping the
    data.
16
17 # %%
18 # Normalizing Data
19 x_train = x_train / 255.0
20 x_test = x_test / 255.0
21
22 # Reshaping Data
23 x_train = x_train.reshape(x_train.shape[0],
    -1)
24 x_test = x_test.reshape(x_test.shape[0], -1)
25
26
27 # %% [markdown]
28 # # Exploratory Data Analysis
29
30 # %%
31 ## How a single entry in dataset looks like
32 instance1 = x_train[2]
33 print(instance1.reshape(1,28,28))
34
35 # %%
36 ## Visualizations
37 colorbar = ['tab:red',
38             'tab:blue',
39             'tab:olive',
40             'tab:orange',
41             'tab:green',
42             'tab:purple',
43             'tab:brown',
44             'tab:pink',
45             'tab:gray',
46             'tab:cyan']
47
48
49 labels = pd.DataFrame(y_train)
50 labels.rename(columns= {0:'Target'}, inplace=
    True)
51 labels.groupby(['Target']).size().plot(kind='
    bar',color=colorbar)
52
53 # %%
54 # visualizing all the 10 categories in the
    dataset
55 plt.figure(figsize=(12,10))
56 x, y = 10, 6
57 for i in range(30):
58     plt.subplot(y, x, i+1)
```

```
59     plt.imshow(x_train[i].reshape(28,28),cmap=
    'magma')
60 plt.show()
61
62 # %%
63 ### Distribution of pixel intensity
64 mean_all = np.mean(x_train)
65 median_all = np.median(x_train)
66 std_all = np.std(x_train)
67
68 print(f"Mean ( training data ): {mean_all:.2f}
    ")
69 print(f"Median (training data): {median_all:.2
    f}")
70 print(f"Standard Deviation (training data): {
    std_all:.2f}")
71
72 # Plot histogram for all images
73 plt.hist(x_train.flatten(), bins=30, color='
    gray', alpha=0.7)
74 plt.title('Pixel Intensity Distribution (
    training data)')
75 plt.xlabel('Pixel Intensity')
76 plt.ylabel('Frequency')
77 plt.show()
78
79
80 # %%
81 ##
82 Tsne = TSNE(n_components=2,perplexity=30,
    max_iter=1000,random_state=24)
83 model_tsne = Tsne.fit_transform(x_train)
84
85 # %%
86 plt.figure(figsize=(12,9))
87 scta_plt = plt.scatter(model_tsne[:,0],
    model_tsne[:,1],
89                         c=y_train.astype(int),
    cmap='tab10', s= 4)
90 plt.colorbar(scta_plt, ticks=np.unique(y_train
    .astype(int)),
91              label="Digits")
92 plt.title('t-distributed stochastic neighbor
    embedding of MNIST')
93 plt.show()
94
95 # %% [markdown]
96 # # Fitting the model
97
98 # %%
99 ##### fitting the model
100 from sklearn.neighbors import
    KNeighborsClassifier
101 KNN_model = KNeighborsClassifier()
102 KNN_model.fit(x_train,y_train)
103
104
105 # %%
106 ## Model prediction
107 predict = KNN_model.predict(x_test)
108
109
110 # %% [markdown]
111 # ## Model evaluation
112
113 # %%
```



```

114
115 from sklearn import metrics
116 evalte = metrics.classification_report(y_test,
117     predict)
118 print(evalte)
119
120 # %%
121 ## vizulaize
122 from sklearn.metrics import confusion_matrix,
123     ConfusionMatrixDisplay
124 conf_mat = confusion_matrix(y_test, predict)
125 conf_mat_disp = ConfusionMatrixDisplay(
126     conf_mat)
127 conf_mat_disp.plot(cmap=plt.cm.BuGn)
128 plt.show()
129
130 # %%
131 from sklearn.metrics import precision_score,
132     recall_score, accuracy_score, f1_score
133 accuracy_score(y_test, predict)
134 precision_score(y_test, predict, average='
135     weighted')
136 recall_score(y_test, predict, average='weighted
137     ')
138 f1_score(y_test, predict, average='weighted')
139
140 # %%
141 ### implement naive bayes Gaussian
142 from sklearn.naive_bayes import GaussianNB
143 model_gnb = GaussianNB()
144 model_gnb_fit= model_gnb.fit(x_train, y_train)
145
146 # %%
147 ## predict
148 y_gnb_pred = model_gnb_fit.predict(x_test)
149
150 # %%
151 gnb_report = metrics.classification_report(
152     y_test, y_gnb_pred)
153 print(gnb_report)
154
155 # %%
156 ### confusion matrix
157 gnb_conf_mat = confusion_matrix(y_test,
158     y_gnb_pred)
159 gnb_conf_mat_disp = ConfusionMatrixDisplay(
160     gnb_conf_mat)
161 gnb_conf_mat_disp.plot(cmap=plt.cm.BuGn)
162 plt.show()
163
164 # %%
165 ## evaluation metrics
166 accuracy_score(y_test, y_gnb_pred)
167 precision_score(y_test, y_gnb_pred, average='
168     weighted')
169 recall_score(y_test, y_gnb_pred, average='
170     weighted')
171 f1_score(y_test, y_gnb_pred, average='weighted')
172
173 # %% [markdown]
174 # ## Auc/ROC
175
176 # %%
177 from sklearn.preprocessing import
178     LabelBinarizer
179 from sklearn.metrics import RocCurveDisplay
180
181 label_binarizer = LabelBinarizer().fit(y_train
182     )
183 y_onehot_test = label_binarizer.transform(
184     y_test)
185 y_onehot_test.shape # (n_samples, n_classes)
186
187 y_score = model_gnb.fit(x_train, y_train).
188     predict_proba(x_test)
189
190 display = RocCurveDisplay.from_predictions(
191     y_onehot_test.ravel(),
192     y_score.ravel(),
193     name="micro-average OvR",
194     color="darkorange",
195     plot_chance_level=True,
196 )
197 _ = display.ax_.set(
198     xlabel="False Positive Rate",
199     ylabel="True Positive Rate",
200     title="Micro-averaged One-vs-Rest\
201         nReceiver Operating Characteristic",
202 )
203
204 # %%
205 from sklearn.preprocessing import
206     LabelBinarizer
207 from sklearn.metrics import RocCurveDisplay
208 import matplotlib.pyplot as plt
209
210 # Binarize the labels
211 label_binarizer = LabelBinarizer().fit(y_train
212     )
213 y_onehot_test = label_binarizer.transform(
214     y_test)
215
216 y_score_knn = KNN_model.fit(x_train, y_train).
217     predict_proba(x_test)
218
219 y_score_gnb = model_gnb.fit(x_train, y_train).
220     predict_proba(x_test)
221
222 # Plot for KNN model
223 display_knn = RocCurveDisplay.from_predictions
224 (
225     y_onehot_test.ravel(),
226     y_score_knn.ravel(),
227     name="KNN micro-average OvR",
228     color="darkorange",
229     plot_chance_level=True,
230 )
231 _ = display_knn.ax_.set(
232     xlabel="False Positive Rate",
233     ylabel="True Positive Rate",
234     title="Micro-averaged One-vs-Rest ROC for
235         KNN and GNB"
236 )
237
238 # Plot for GNB model

```

```
225 RocCurveDisplay.from_predictions(  
226     y_onehot_test.ravel(),  
227     y_score_gnb.ravel(),  
228     name="GNB micro-average OvR",  
229     color="blue",  
230     ax=display_knn.ax_,  
231     plot_chance_level=False,  
232 )  
233  
234 # Show the plot  
235 plt.show()
```